

Distributed Monolith: Challenges and strategies in breaking Databases for Microservices

MAREK MAŠKO



16 edycja konferencji SQLDay

13-15 maja 2024, WROCŁAW + ONLINE



partner platynowy

lingaro

partner złoty



partner srebrny



About the Author



MAREK MAŚKO

Sr Engineering Manager
CAE Flight Services Poland

E-MAIL

MarekMasko@cae.com


LINKEDIN

linkedin.com/in/MarekMasko/

TWITTER

@MarekMasko

Microservices



An unrelenting trend



Source: trends.google.com

What are microservices?



”

Microservices are an architectural and organizational approach to software development where software is composed of **small independent services** that **communicate over well-defined APIs**. These services are owned by small, self-contained teams.



Microservices architectures make applications **easier to scale** and **faster to develop**, enabling innovation and accelerating time-to-market for new features.

”

Source: aws.amazon.com

What are microservices?



”

Microservices are an architectural approach to building applications where **each core function, or service, is built and deployed independently.**



Microservice architecture is **distributed** and **loosely coupled**, so one component's failure won't break the whole app. **Independent components** work together and communicate with well-defined API contracts.

”

Source: azure.microsoft.com

What are microservices?



”

Microservices (...) refers to an architectural style for developing applications. Microservices allow a large application to be separated into **smaller independent parts**, with each part **having its own realm of responsibility**.



To serve a single user request, a microservices-based application can call on many internal microservices to compose its response.

”

Source: cloud.google.com

What are microservices?



”

(...) **microservice** architectural style is an approach to developing a single application as a **suite of small services**, each running in its own process and **communicating with lightweight mechanisms**, often an HTTP resource API.



These services are **built around business capabilities** and **independently deployable** by fully automated deployment machinery. There is a **bare minimum of centralized management** of these services, which may be written in different programming languages and use different data storage technologies.

”

Source: martinfowler.com

Small
Independent
Having own responsibility
Loosely coupled

What problems do they solve?



What problems do **you want to** solve?



Scalability?

Deliver faster and more often?

Too big codebase?

Embrace the new technology?

Are microservices a good choice for me?




I don't know...

...and we won't talk about that!

Application migration patterns


- Strangler Fig Application
- UI Composition
- Branch By Abstraction
- Parallel Run
- Decorating Collaborator
- Change Data Capture
- ...



**Where is the database
in all of this?**



Databases are different...



...databases have state!

Microservice architecture style...



Microservice architecture style...



Microservices

Microservice architecture style...



Microservices \Rightarrow Stateless apps

Microservice architecture style...



Microservices \Rightarrow Stateless apps

State


Microservice architecture style...



Microservices \Rightarrow **Stateless** apps

State \Rightarrow **Data services**

It's just a shift of the **problem**



Distributed monolith

Decomposition of a database



You **can** put it off initially

But it **can't** be delayed forever

Microservice characteristics

- Componentization via services
- Organized around business capabilities
- Products not projects
- Smart endpoints and dumb pipes
- Decentralized governance
- **Decentralized data management**
- Infrastructure automation
- Design for failure
- Evolutionary design

Microservices should own their own data






”

One of the things I see people having the hardest time with is the idea that **microservices should not share databases**.

If one service wants to access data held by another service, then it should go and ask that service for the data it needs.

”

Source: Book „Monolith to Microservices” by Sam Newman



Big bang or smooth migration?

App and DB at the **same time**

vs

Baby steps...

Zero-downtime migrations



Prerequisites

Treat DB code same as an App code:

- Version control system
- Versioning
- Automated deployment

No data loss allowed!

Change column name



```
1  -- ALTER TABLE dbo.customers RENAME COLUMN wrong TO correct;
2  EXEC sp_rename 'dbo.Customers.Wrong', 'Correct', 'COLUMN';
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
```

Change column name



```
1 ALTER TABLE dbo.Customers ADD Correct NVARCHAR(60);
2 GO
3
4 UPDATE dbo.Customers SET Correct = Wrong WHERE CustomerId BETWEEN 1 AND 100;
5 UPDATE dbo.Customers SET Correct = Wrong WHERE CustomerId BETWEEN 101 AND 200;
6 UPDATE dbo.Customers SET Correct = Wrong WHERE CustomerId BETWEEN 201 AND 300;
7 UPDATE dbo.Customers SET Correct = Wrong WHERE CustomerId BETWEEN 301 AND 400;
8 UPDATE dbo.Customers SET Correct = Wrong WHERE CustomerId BETWEEN 401 AND 500;
9 UPDATE dbo.Customers SET Correct = Wrong WHERE CustomerId BETWEEN 501 AND 600;
10 UPDATE dbo.Customers SET Correct = Wrong WHERE CustomerId BETWEEN 601 AND 700;
11
12 ALTER TABLE dbo.Customers ALTER COLUMN Correct NVARCHAR(60) NOT NULL;
13 ALTER TABLE dbo.Customers DROP COLUMN Wrong;
14
15
16
17
```

Remember



Rehearse your migrations

and

Check your data
between migrations steps


A few things to consider




Logical vs Physical decomposition




Logical vs Physical decomposition





Logical vs Physical decomposition




Logical vs Physical decomposition




Logical vs Physical decomposition



Logical vs Physical decomposition





Logical decomposition



Physical decomposition

Referential Integrity




Move FK relationship to the App code

Move from Strong Consistency to Eventual Consistency




Move from Strong Consistency to Eventual Consistency




Strong Consistency

Move from Strong Consistency to Eventual Consistency

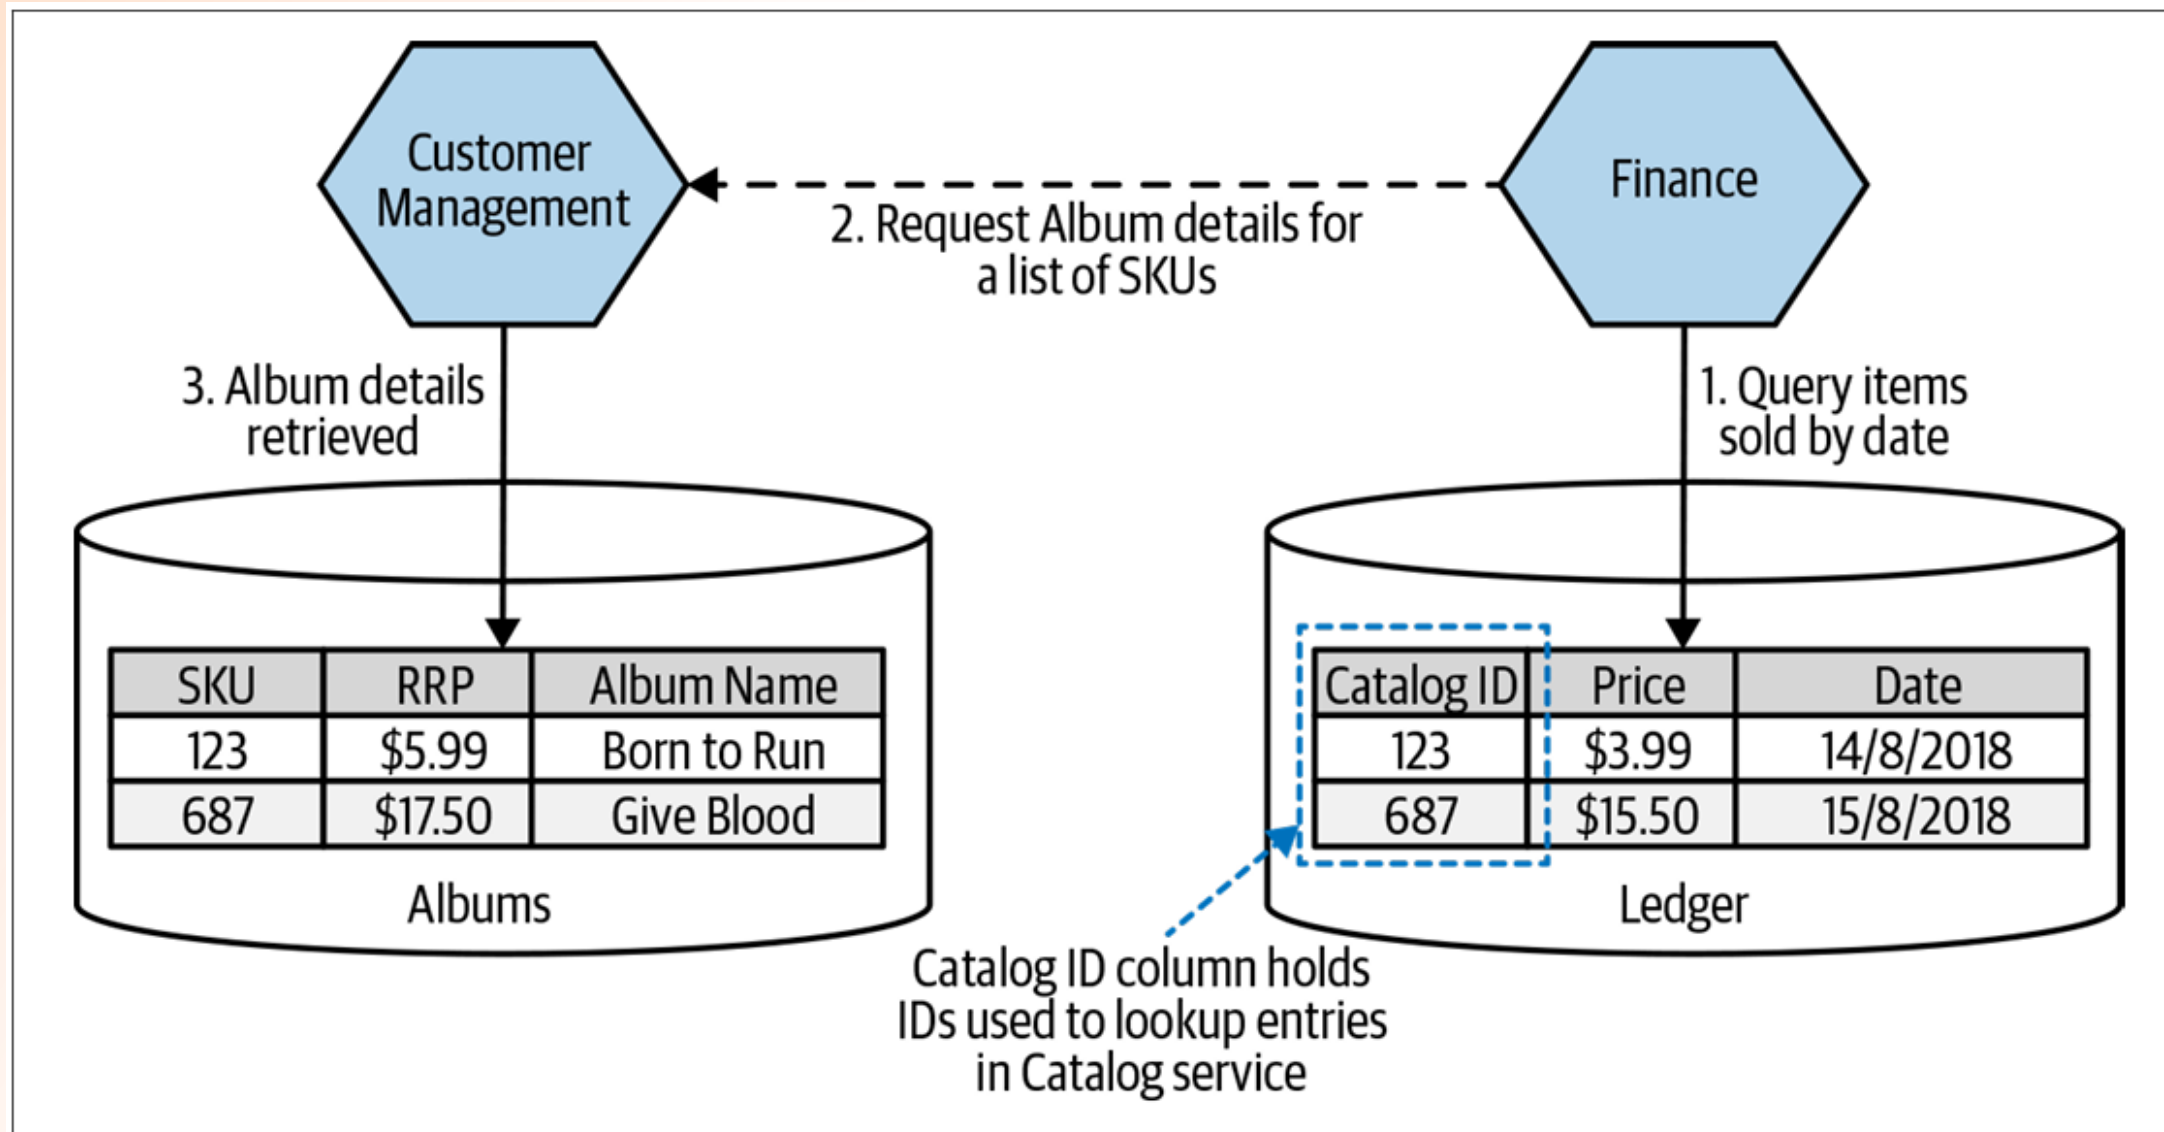


Strong Consistency



Eventual Consistency

Data consistency – Deleting data



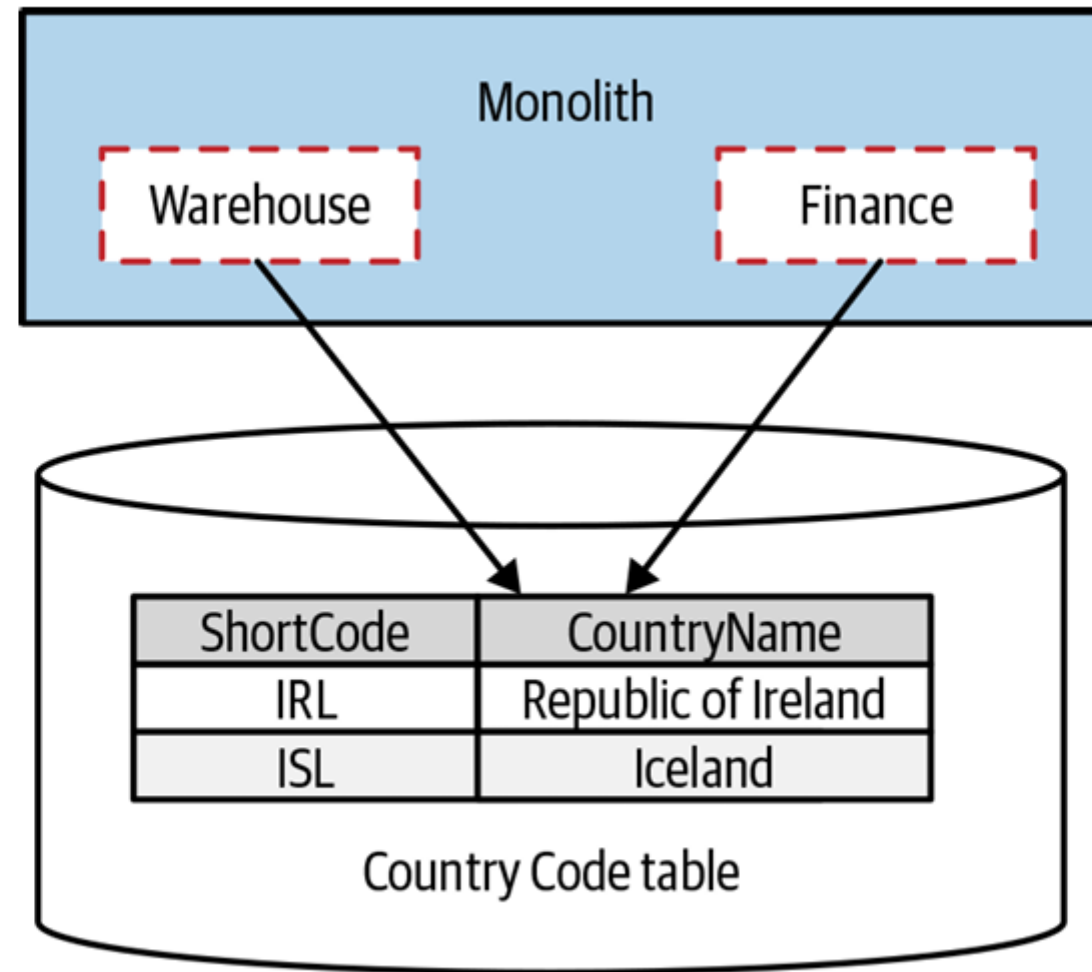
Data consistency – Deleting data

- Check before deletion
 - Stop new references being created – locking
 - Reverse dependencies
- Handle deletion gracefully
 - „Information not available”
 - Not found vs Gone
 - Subscribe to deletion event and copy data locally
- Don't allow deletion
 - Soft delete


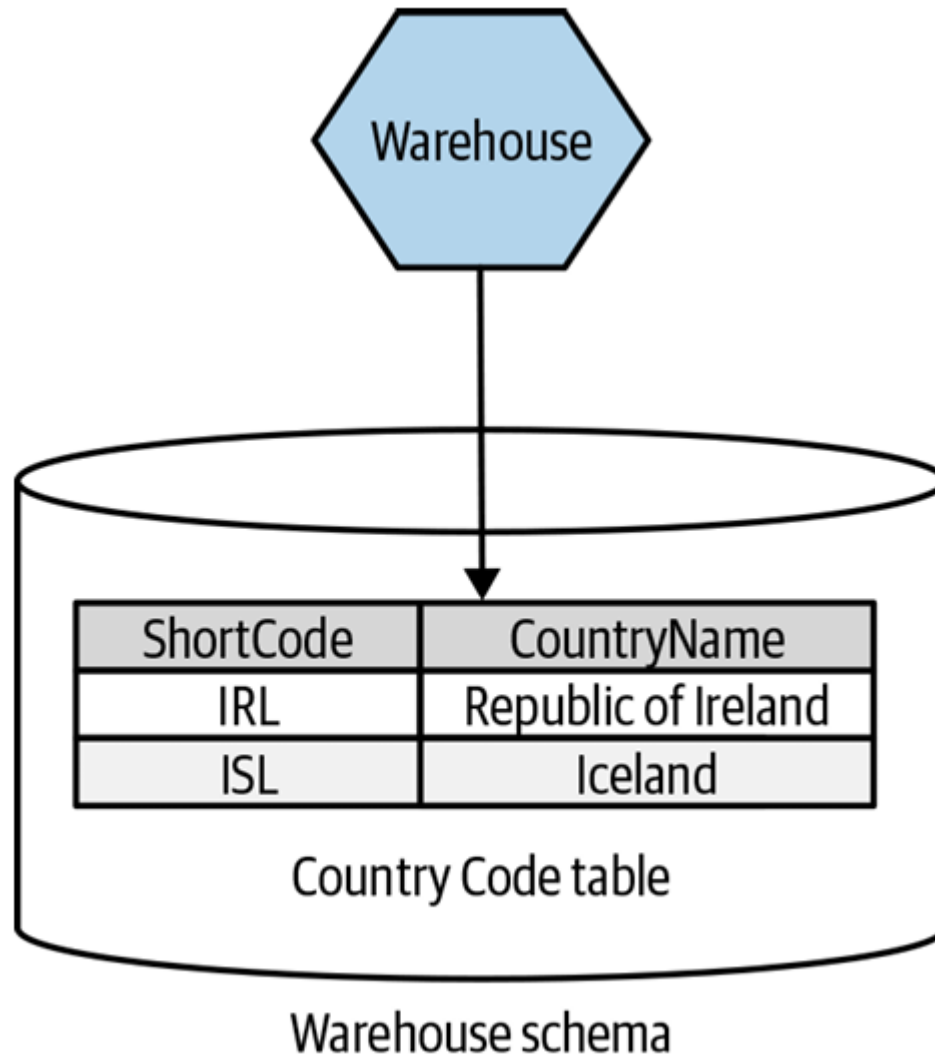
**What about recovering the DB
to an earlier state?!**

**We created an issue that doesn't
exists in the monolithic system**

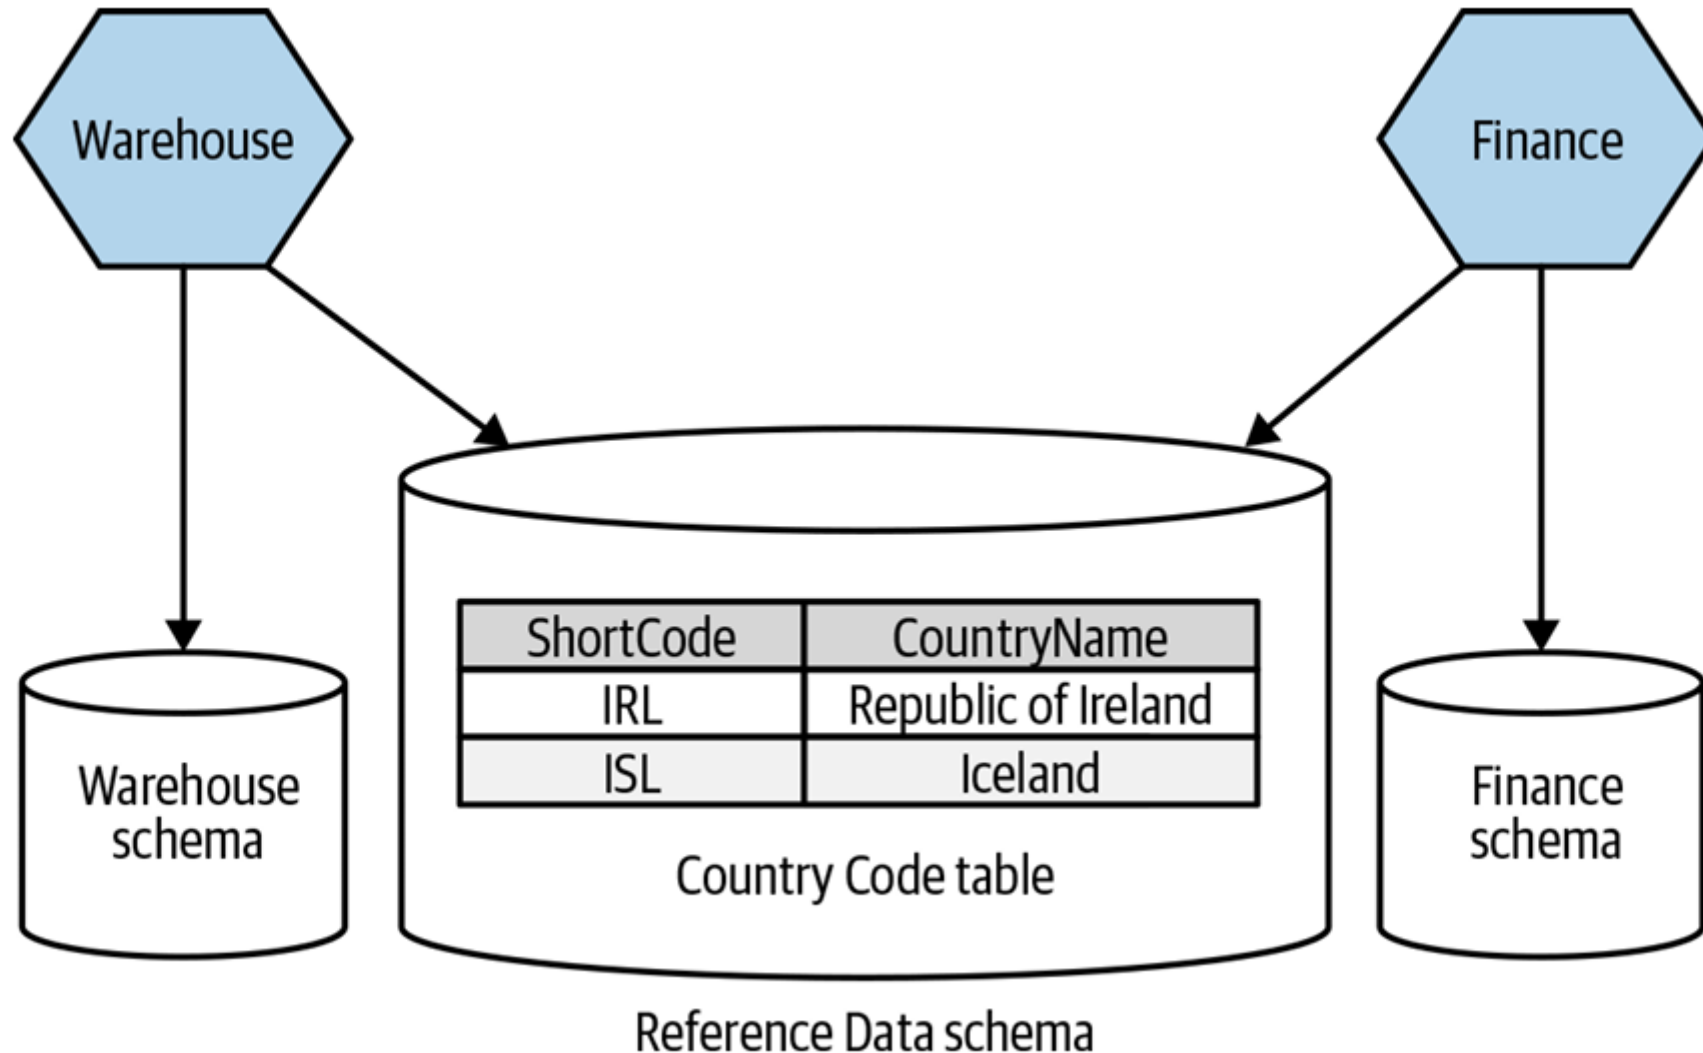
Shared static data



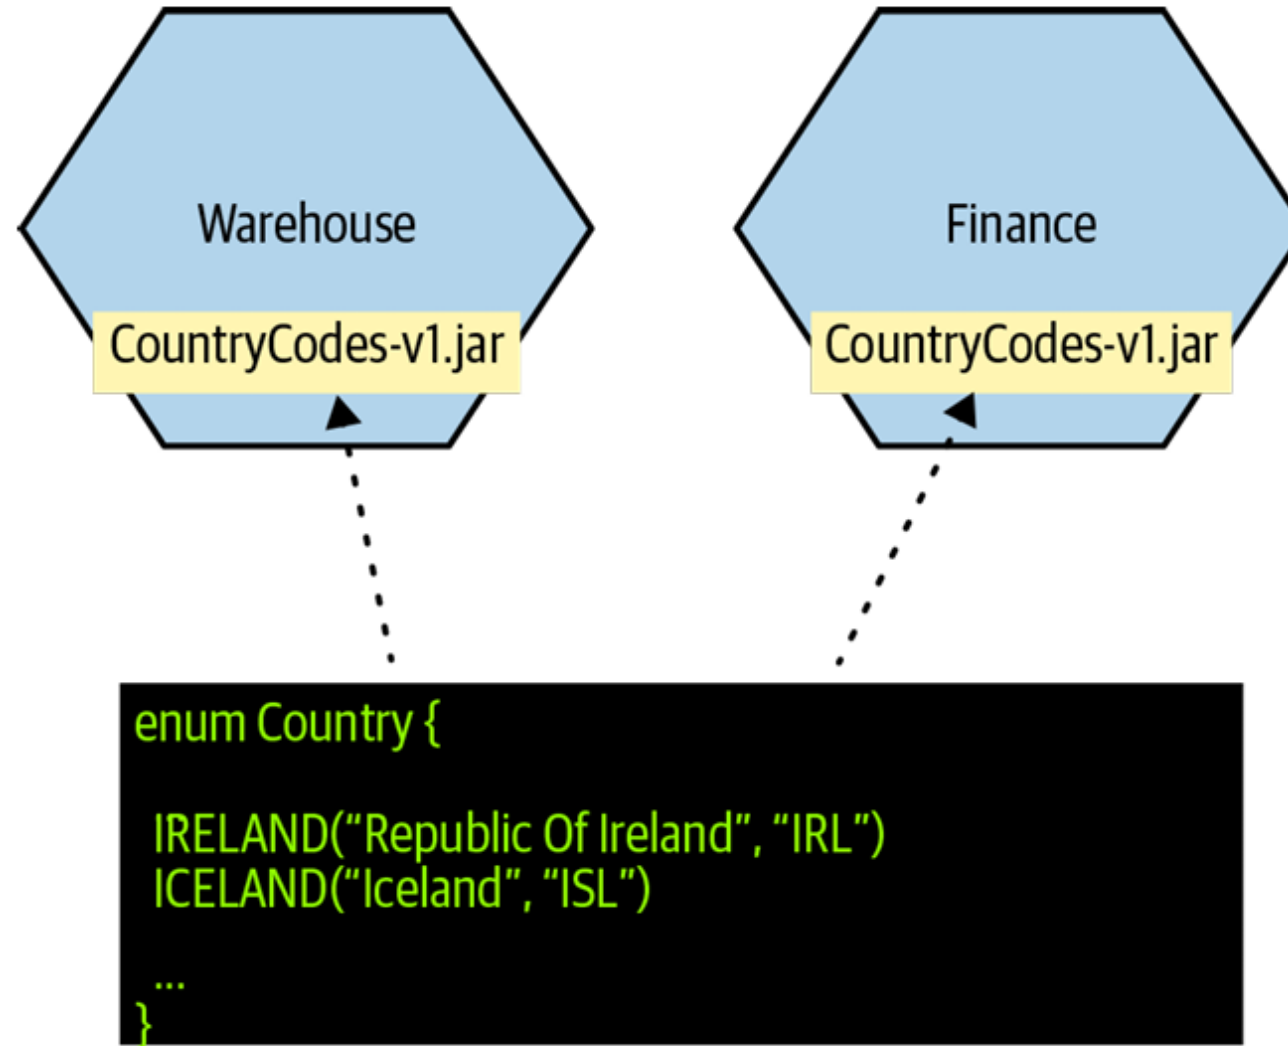
Pattern: Duplicate static reference data




Pattern: Dedicated reference data scheme



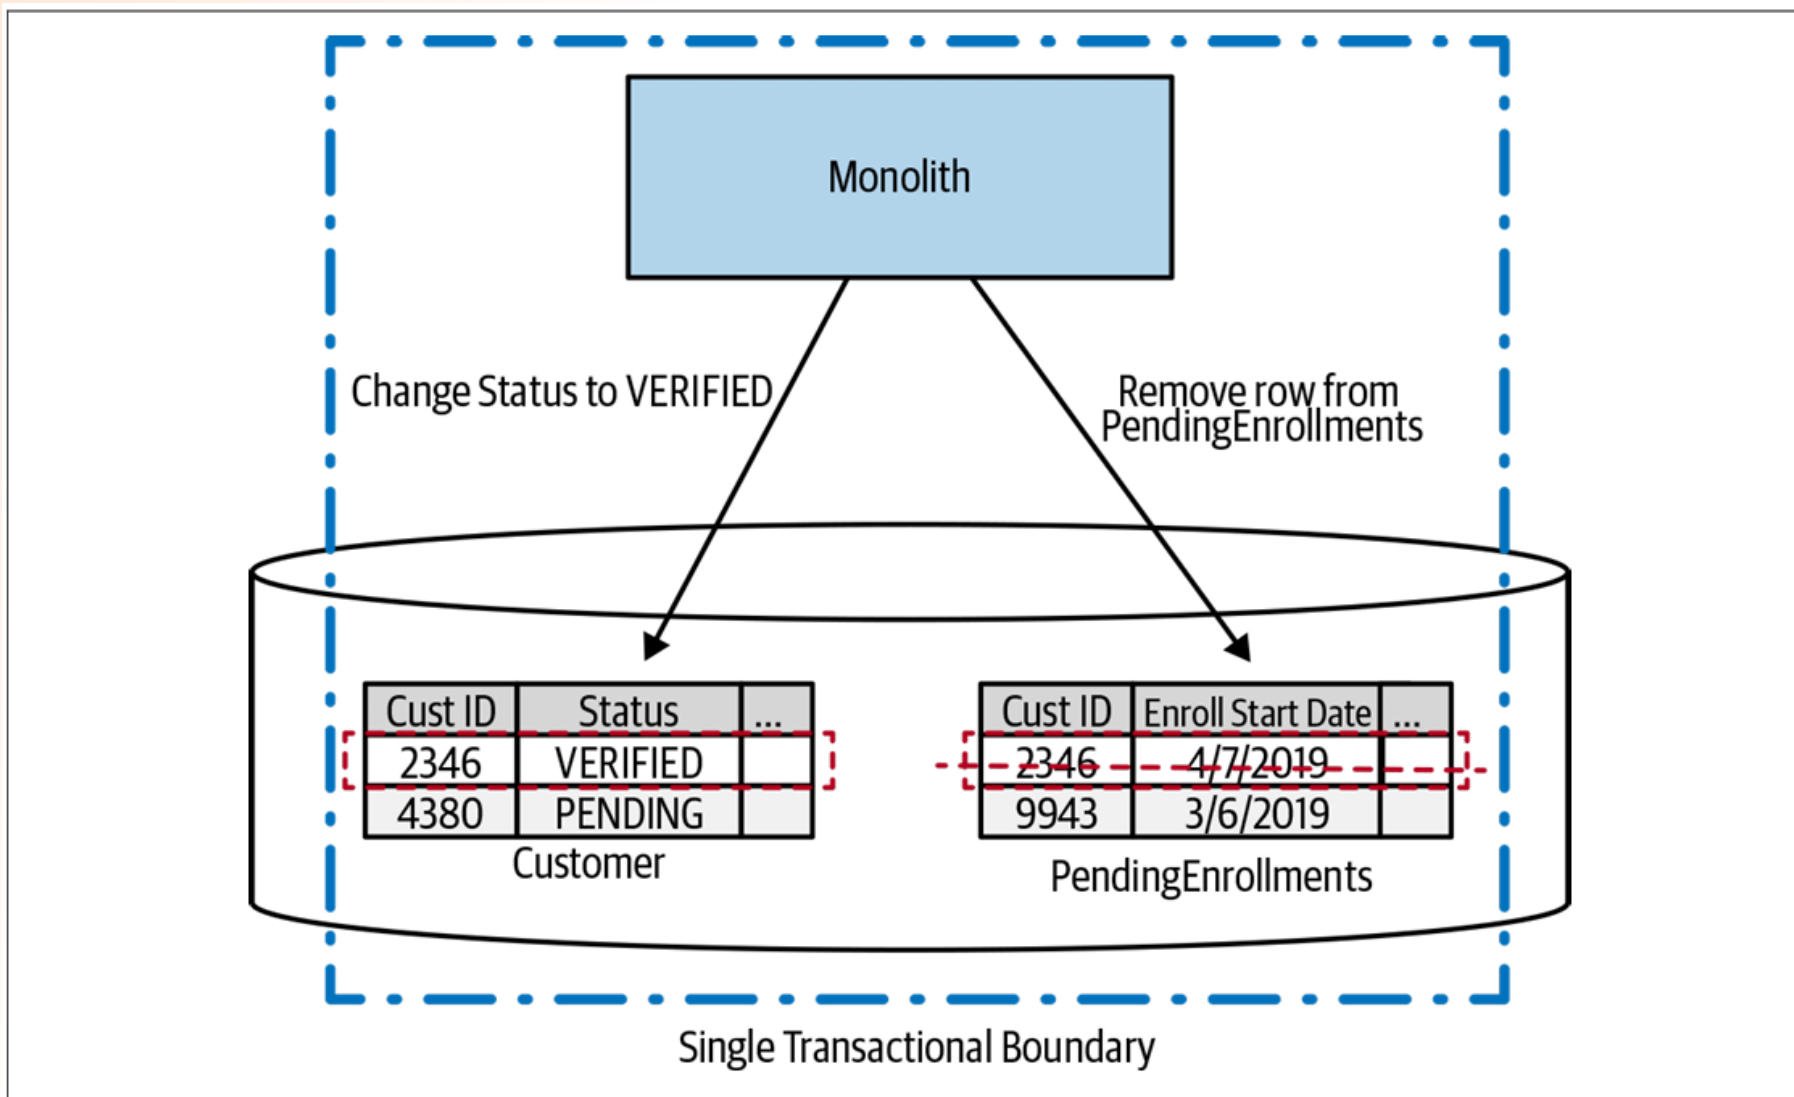
Pattern: Static reference data library



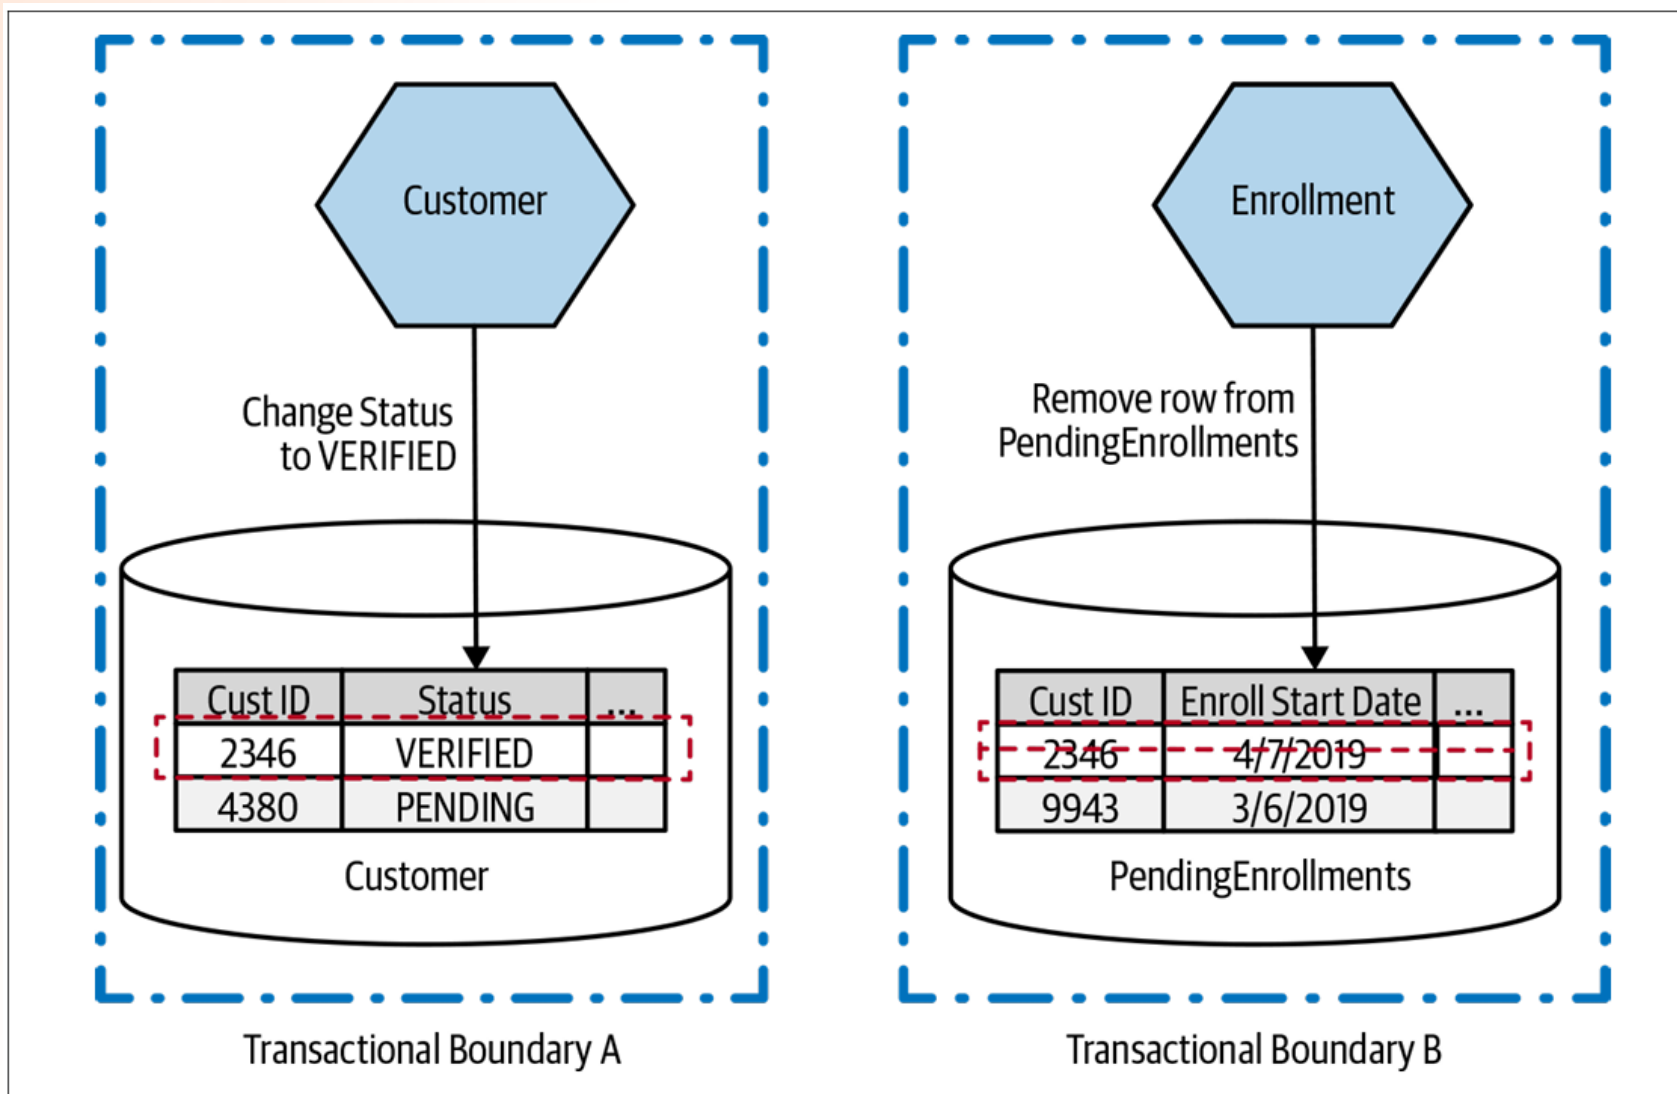
Pattern: Static reference data service



Transactions

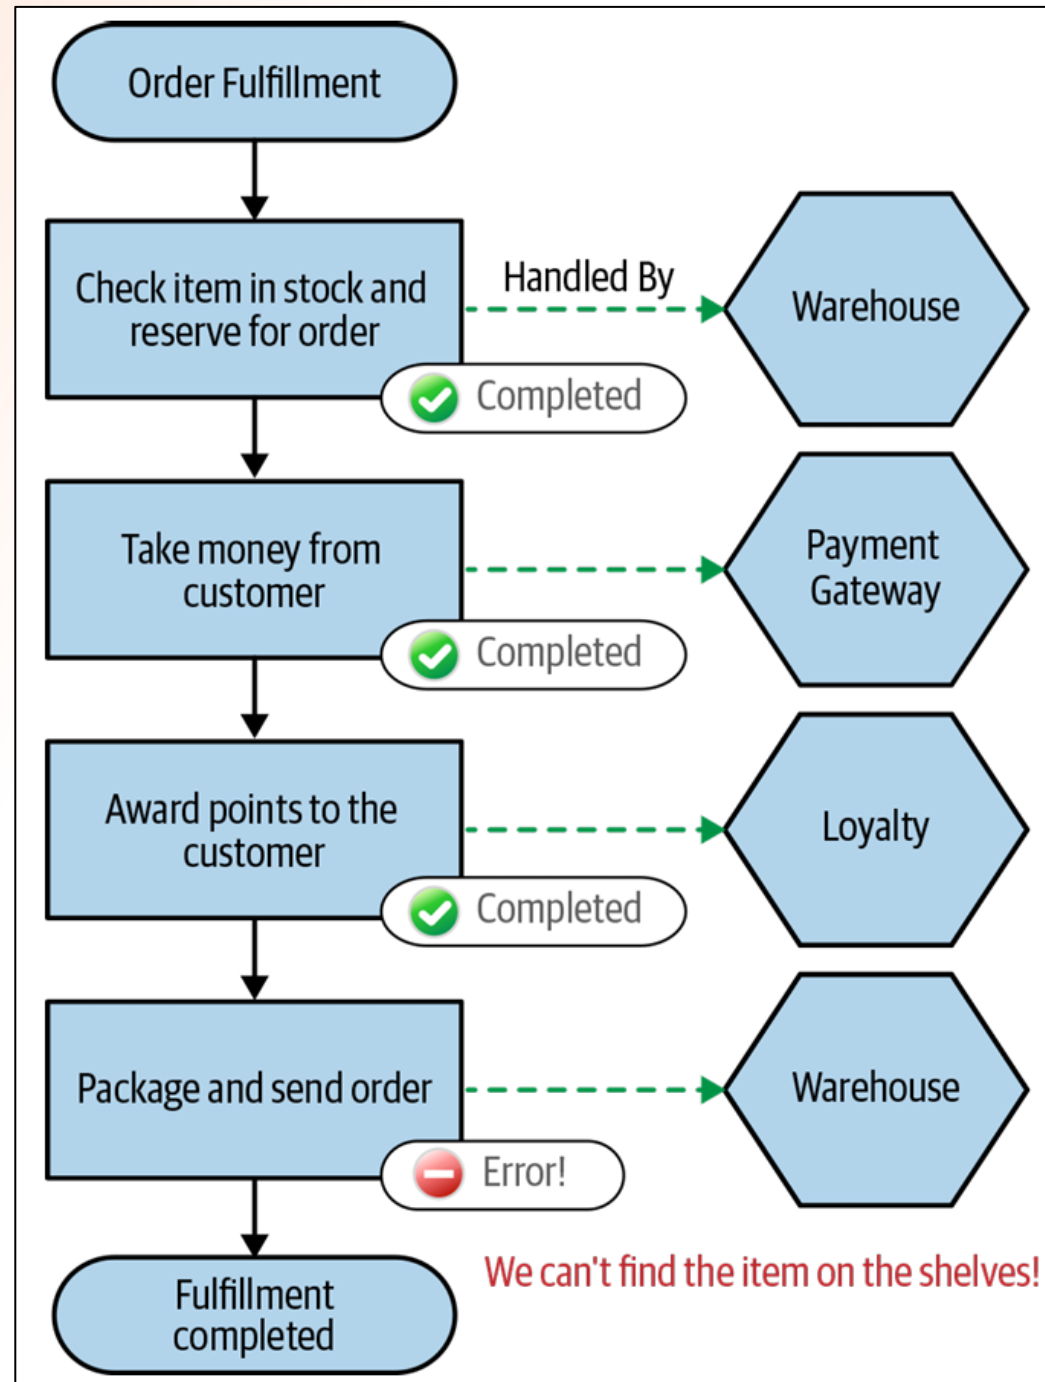


Transactions



Lack of atomicity is a serious issue for systems that previously relied on this property




Sagas






Change the database type?








Change the database type?



Change the database type?




Change the database type?

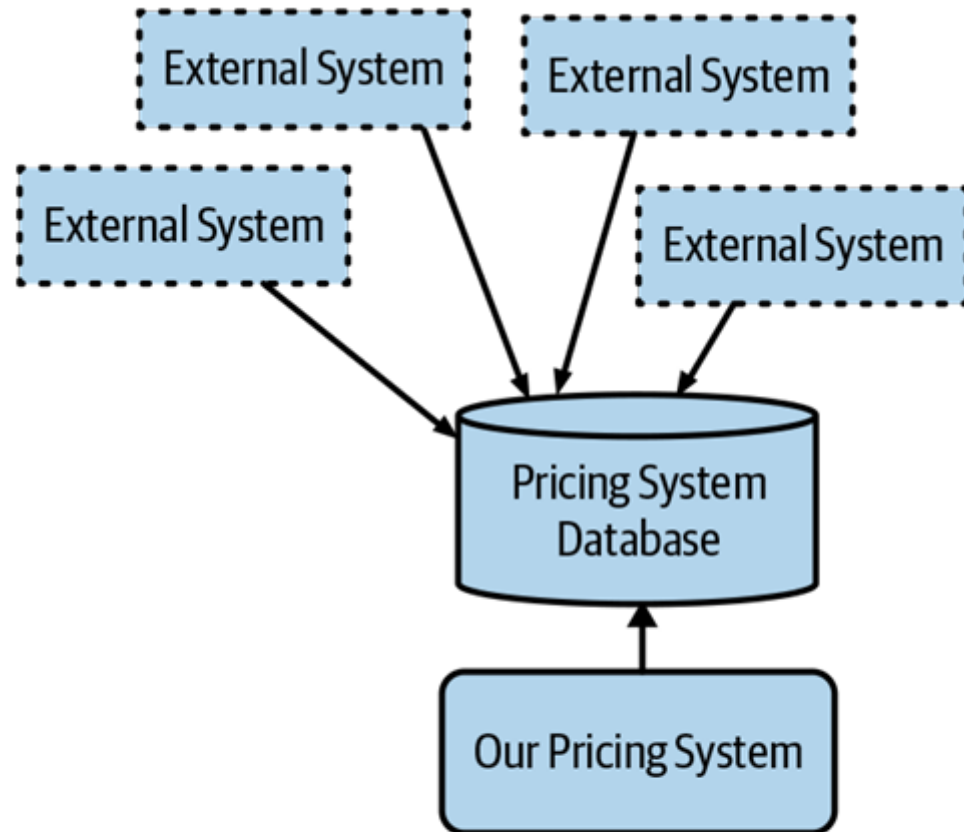


Migration Patterns

Expose Data




Database views

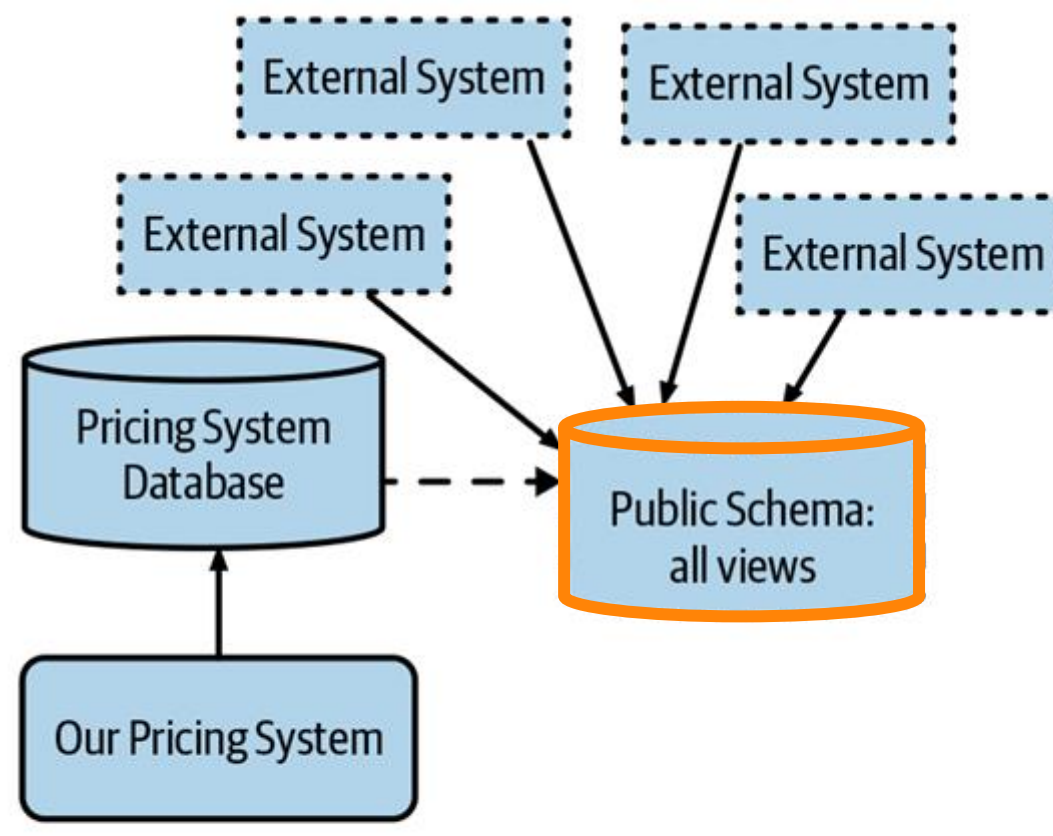


Before: External systems have direct DB access

Database views



Before: External systems have direct DB access



After: External systems are redirected to read data from views, allowing the schema for the pricing system to change

Database views

Considerations:

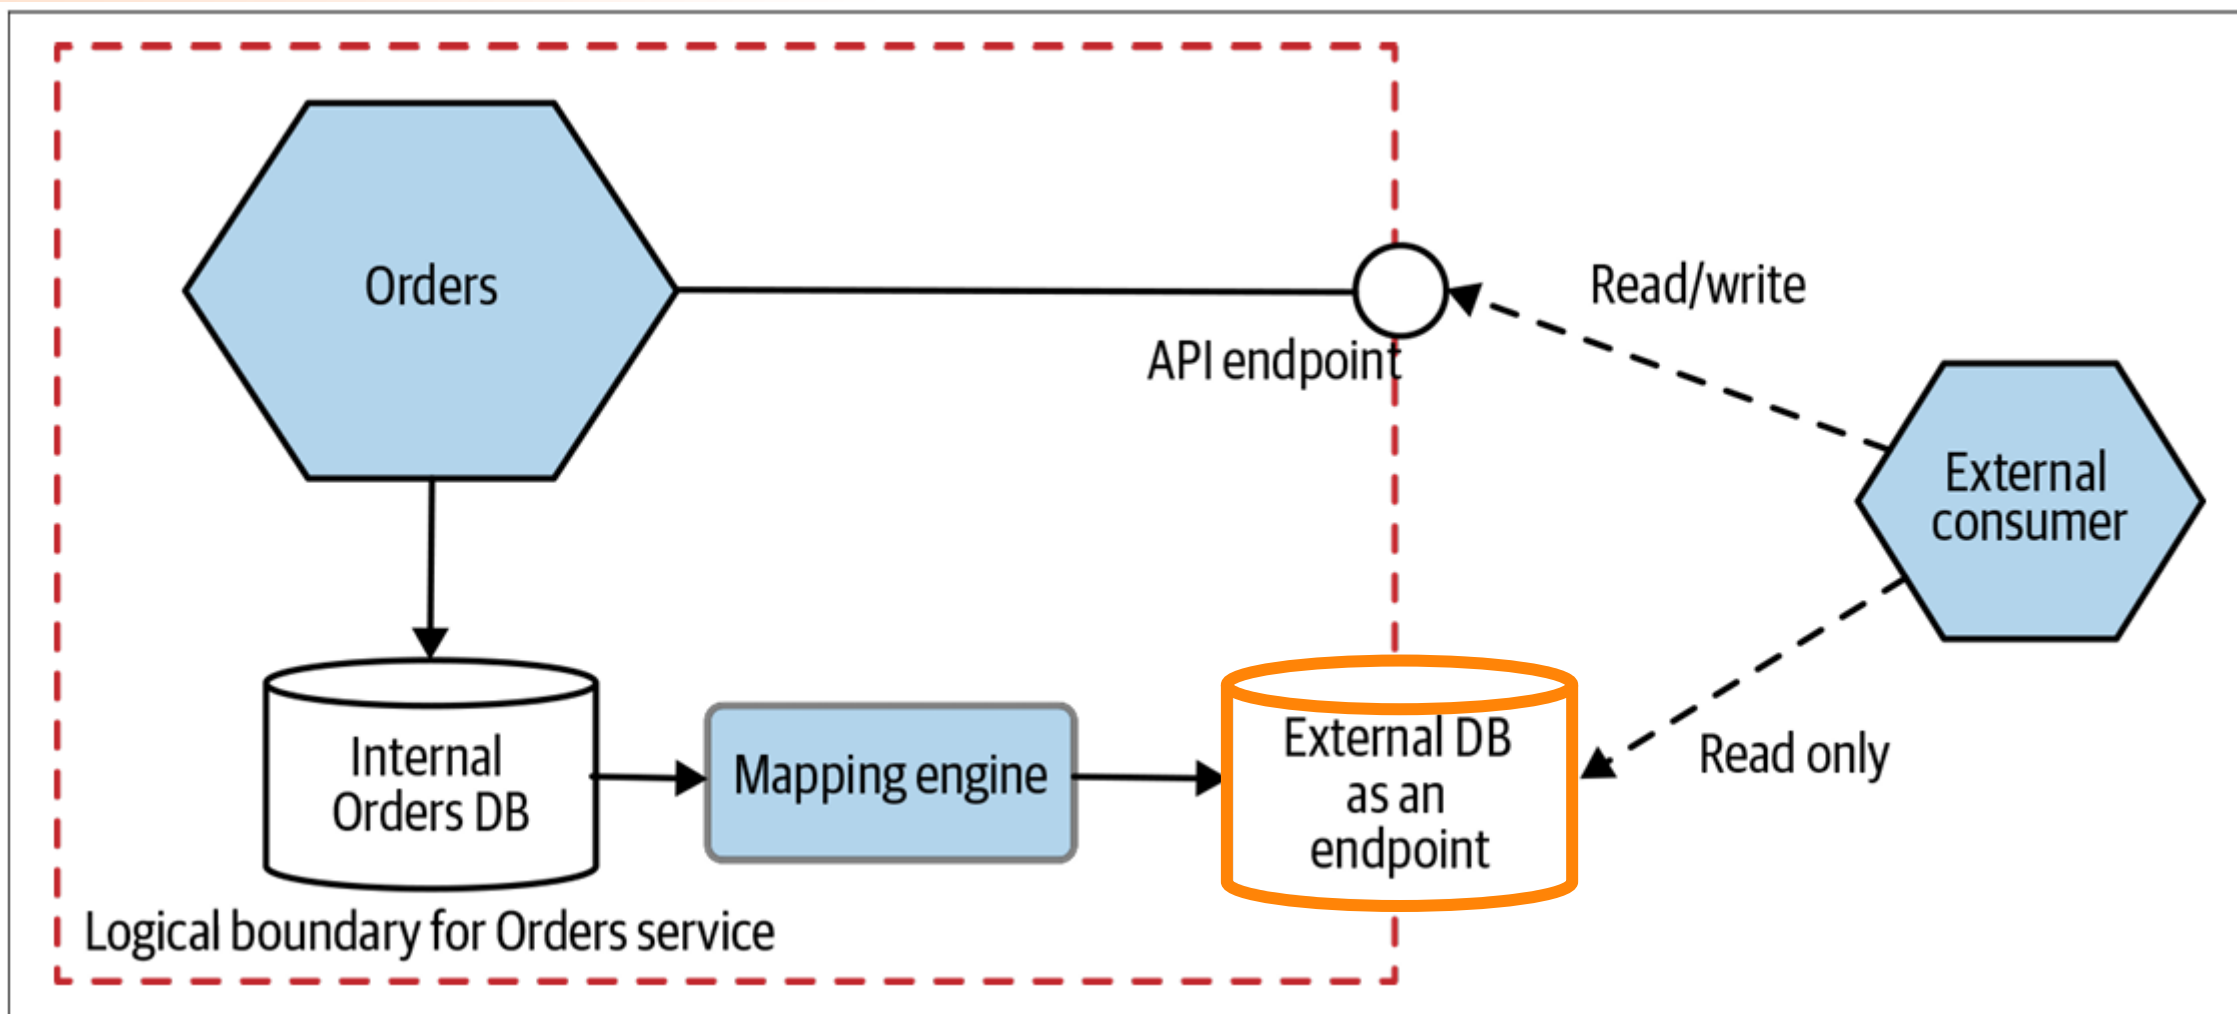
- Easiest strategy to implement
- Largest support from DBMS vendors
- Possible performance issues
- Strong consistency
- One database must be reachable by the other
- Updatable depending on DBMS support

Database materialized views

Considerations:

- Better performance
- Strong or eventual consistency
- One database must be reachable by the other
- Updatable depending on DBMS support

Database as an interface




Database as an interface

Considerations:


- Requires mapping engine
- Strong or eventual consistency
- Read-only integration
- Can be implemented in many ways

Migration Patterns

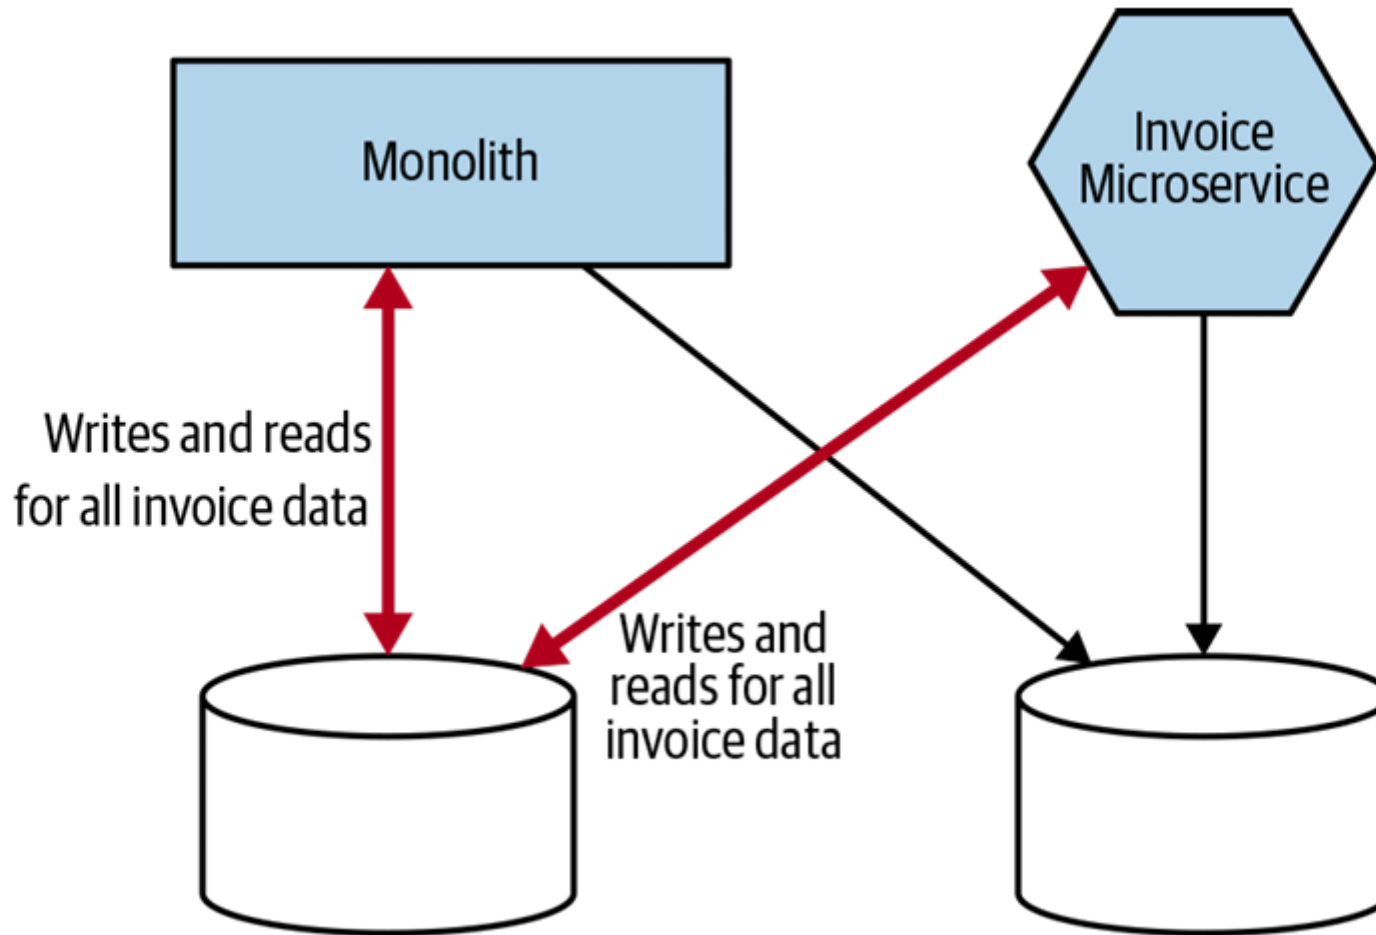
Synchronize Data



Synchronize data in the app



Synchronize data in the app

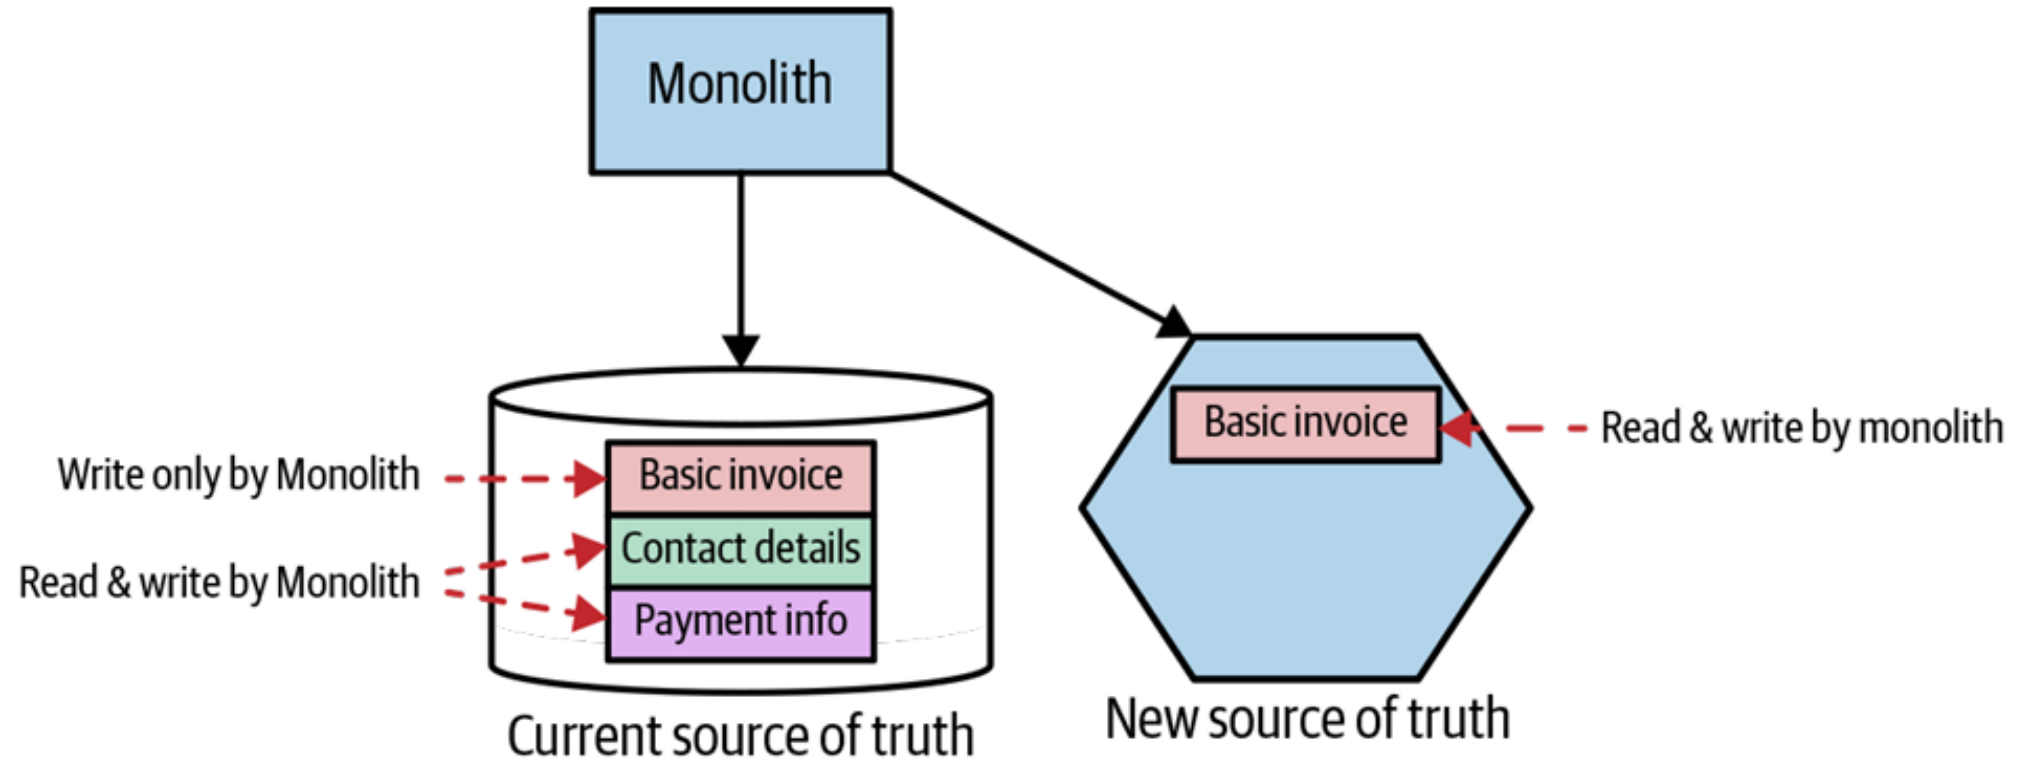


Synchronize data in the app

Considerations:

- Double write logic in the app
- Strong consistency
- High cohesion and low coupling


Tracer Write



Write - inconsistency

- Write to one source
- Send writes to both sources
- Send writes to either source

How to keep this in sync?



Database trigger

Considerations:

- Depends on the DBMS support
- Strong consistency
- One database must be reachable by the other



Transactional code

Considerations:

- Any code: usually stored procedures or distributed transactions
- Strong consistency
- Possible cohesion/coupling issues
- Possible performance issues
- Updatable depending on how it is implemented



ETL - Extract Transform Load

Considerations:

- Many available tools
- Requires external trigger (usually time-based)
- Can aggregate from multiple data sources
- Eventual consistency
- Read-only integration



Replication

Considerations:

- A few available options
- The same source and target schemas
- Strong or Eventual consistency
- Updatable depending on replication type




Change Data Capture

It's a process of identifying and capturing changes made to data in a database and then delivering those changes in real time to a downstream process or system.



Change Data Capture - Debezium




Change Data Capture

Considerations:

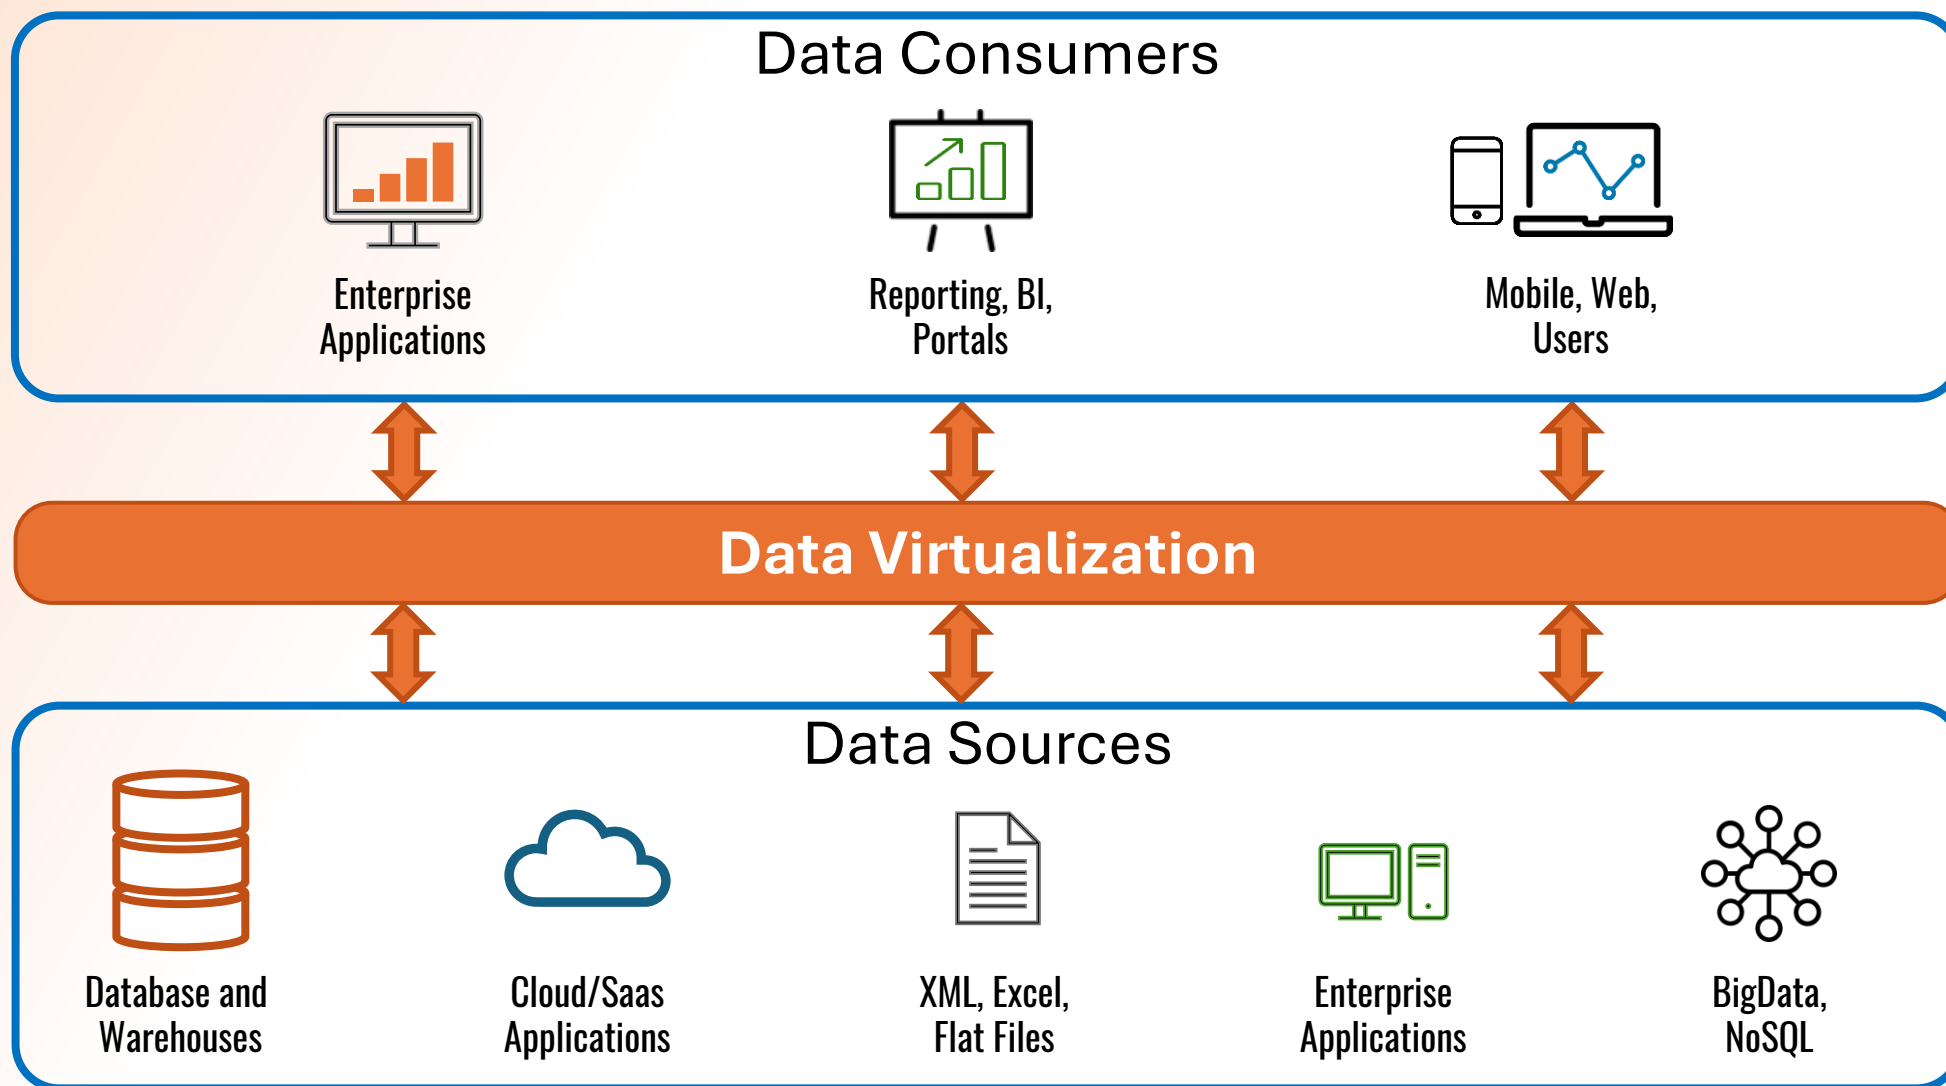
- Read data source is updated through a stream of events
- Eventual consistency
- Usually combined with a message bus – Kafka?
- High scalability

Migration Patterns

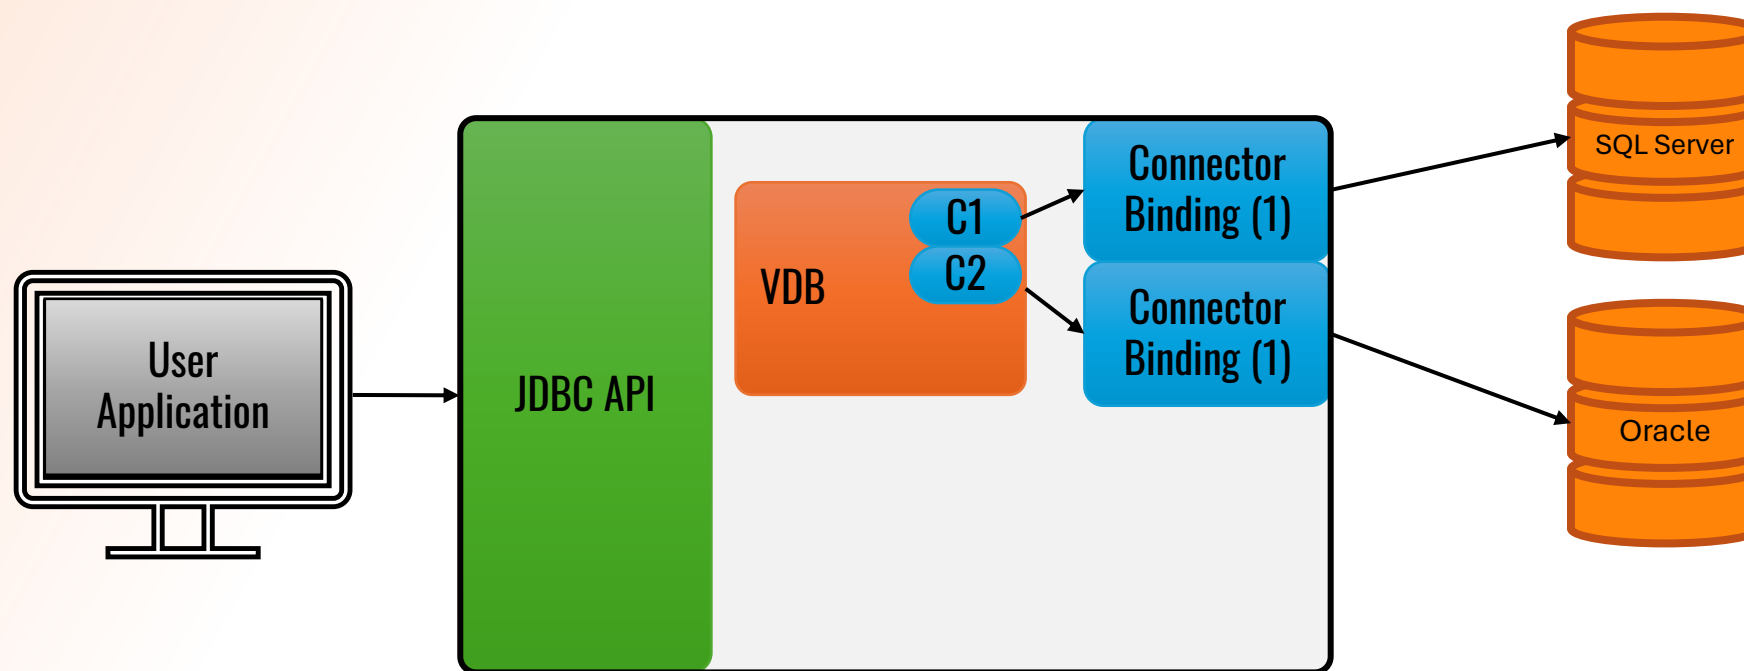
Virtualize Data



Data virtualization



Data virtualization



Data virtualization

Many tools available:

- Teiid.io
- Denodo
- ...
- SQL Server - Polybase

Data virtualization

Considerations:

- Real-time access option
- Strong or eventual consistency
- Can aggregate from multiple data sources
- Updatable depending on data virtualization platform
- Additional layer can introduce latency

Summary



”

Code is easy;
State is hard.

”

-Edson Yanaga

Data integrity

Data consistency

Transaction boundaries

Data mapping engine

Be prepared...



...for **issues that didn't exist in**
the monolithic system!

Remember





Rehearse your migrations

and

Check your data
between migrations steps

Resources





Data
Community



16 edycja konferencji SQLDay

13-15 maja 2024, WROCŁAW + ONLINE



partner platynowy



partner złoty



partner srebrny

